

NORTHSTAR FINANCIAL SERVICES

AI GOVERNANCE & MODEL SAFETY FRAMEWORK V2.1

DOCUMENT ID	EFFECTIVE DATE	POLICY OWNER	CLASSIFICATION
POL-AI-2025-01	Nov 01, 2025	AI Governance Board & CISO	Internal Restricted

1. PURPOSE & REGULATORY SCOPE

As Northstar Financial Services (NFS) integrates artificial intelligence, large language models (LLMs), and automated deep agent systems into core operations, this framework establishes mandatory safety, transparency, ethical, and regulatory baselines. All third-party AI platform vendors must satisfy these rules before deployment in customer-facing or financial decision-making workflows.

2. MODEL TRANSPARENCY & EXPLAINABILITY STANDARDS

Black-box models operating without explainable reasoning mechanisms are restricted from automated financial approvals. Vendors must provide comprehensive model disclosures:

- Model Cards:** Vendors must submit detailed Model Cards outlining architecture size, training cutoff dates, foundational datasets, known failure modes, and benchmark performance metrics.
- Provenance & Training Data Disclosure:** Vendors must certify that training corpora did not incorporate copyrighted financial records, proprietary NFS trade secrets, or illicitly harvested PII.
- Explainable AI (XAI):** For credit scoring, automated trading recommendations, or fraud flagging, models must support deterministic feature attribution (e.g., SHAP/LIME values or structured citation trees) so human officers can audit decisions.

3. DATA USAGE, PRIVACY & ANTI-CONTAMINATION MANDATES

Client confidentiality and data isolation are paramount when interacting with third-party foundational models:

GOVERNANCE AREA	NFS MANDATORY BASELINE STANDARD	VERIFICATION MECHANISM
Zero Data Retention	Vendor API endpoints must enforce zero data retention (ZDR). Inputs, prompts, and inference outputs must not be persisted beyond active session computation.	Data processing agreement (DPA) review & API logging audit.
No Secondary Training	Explicit contractual prohibition against using NFS prompts, fine-tuning data, or telemetry to train, retrain, or align public or foundational models.	Legal clause inspection & architectural guarantee.
Hallucination & Drift	Models deployed in customer workflows must maintain a factual hallucination rate below 0.05% on standardized financial benchmark testing suites.	Automated red-teaming & evaluation benchmark reports.
Guardrails & Safety	Mandatory input/output moderation filters to block prompt injection, jailbreaks, toxic outputs, and unauthorized data exfiltration.	Security penetration test & adversarial prompt stress testing.

AI Governance Rule: Any AI vendor that reserves the right to review customer prompts or outputs for "product improvement" or "quality assurance via human review" is automatically disqualified for Tier-1 financial workflows.

4. HUMAN-IN-THE-LOOP (HITL) & DECISION AUTONOMY TIERS

NFS enforces strict operational boundaries regarding autonomous AI actions. AI systems are categorized by autonomy thresholds:

AUTONOMY TIER	PERMITTED OPERATIONAL SCOPE	HUMAN OVERSIGHT REQUIREMENT
Tier A: Autonomous	Routine document parsing, data formatting, vector indexing, syntax checks, and automated initial policy retrieval.	Zero human intervention required; asynchronous spot-audit logging.
Tier B: Assisted	Drafting vendor evaluation scorecards, synthesizing audit summaries, and proposing preliminary risk classifications.	Mandatory human review and sign-off before official report generation.
Tier C: Restricted	Automated credit decisions, loan pricing, trade execution, final vendor approval/rejection, and legal contract redlining.	Strict human-in-the-loop (HITL) mandatory approval. AI acts solely as recommender.

5. AI GOVERNANCE RED-LINES & MANDATORY VETO CRITERIA

The AI Governance Agent will issue an immediate veto (*Reject Status*) if a vendor proposal or system architecture violates any of the following mandatory red-lines:

VETO Training on NFS Data

Vendor claims ownership of generated outputs or reserves rights to use NFS query inputs and documents for model training.

VETO Unexplainable Black-Box Scoring

Providing automated credit or compliance rejection scores without traceable feature attribution or evidence citations.

VETO Unmoderated Open-Weights

Deploying unaligned open-weight models lacking enterprise safety filters, jailbreak prevention, and toxicity guardrails.

VETO Third-Party Data Reselling

Sharing metadata, usage volume, or financial workflow patterns with external data brokers or parent enterprise partners.

6. CONTINUOUS MODEL MONITORING & AUDIT TRAIL

Approved AI vendors must integrate with the NFS Audit Logger MCP server:

- Immutable Logging:** Every prompt hash, tool invocation, token count, and model response must be logged to an append-only audit trail for regulatory inspection.
- Bias & Fairness Audits:** Quarterly algorithmic bias evaluations must be submitted to the AI Governance Board to ensure equitable processing across protected customer demographics.